

Teaching founders to spot and verify AI fabrications

"Your pitch deck has numbers in it. If your AI made them up, investors will destroy you. Let's build a BS detector."

Overview

Students learn that AI models confidently generate false information — and how to build a verification layer into their everyday business workflows using free tools they already have.

- Duration: 25 minutes
- Tools: llm CLI (Groq free tier — no credit card)
- Setup: 5 minutes (already installed from Week 1)
- Format: Pairs, hands-on with terminals
- Fits: Week 2 or 3, alongside The Translator or Spy Agency

Learning Objectives

- Understand what AI hallucination is and why it happens
- Experience firsthand how confidently wrong an AI can be
- Build a two-prompt verification workflow
- Apply skepticism to AI outputs in business contexts
- Connect AI literacy to real founder risk (pitch decks, research, customer claims)

Activity Flow (25 min)

Step 1 — The Trap (5 min)

Each pair picks a business claim (provided or their own). Examples:

- "70% of startups fail in year one"
- "The average small business spends \$5,000/month on marketing"
- "90% of consumers trust online reviews as much as personal recommendations"

Step 2 — First Prompt (5 min)

Students send the claim to the AI asking: "Is this true? Cite your sources."

The AI responds with a confident, plausible-sounding answer — often with fabricated citations.

Step 3 — The Reveal (5 min)

Pairs compare answers. The facilitator reveals: the AI just made it up. No sources exist.

Discussion: Why does the AI sound so sure when it's wrong?

Step 4 — The Verifier (5 min)

Students run a second prompt that cross-checks the first answer. They compare side by side.

Step 5 — Debrief (5 min)

What did this teach you about using AI for business research? How would you protect your pitch deck?

The Two Prompts

PROMPT 1: The Claim	PROMPT 2: The Verifier
You are a business research assistant. I need to verify this claim for a pitch deck: "70% of startups fail in their first year." Is this accurate? Please cite specific sources and studies.	You are a fact-checker. Someone claims that 70% of startups fail in their first year. What does actual data from reliable sources (SBA, CB Insights, Bureau of Labor Statistics) say? Flag any common misconceptions.

Run them back-to-back and compare the answers.

CLI Commands

Set your key once:

```
$ llm keys set groq
```

```
$ llm -m groq "You are a business researcher. Is it true that 70% of startups fail in year one? Cite sources."
```

Then the verifier:

```
$ llm -m groq "You are a fact-checker. What does actual data from SBA and BLS say about startup failure rates?"
```

Why This Works

- Entrepreneurship-first: Opens with a real founder problem (pitch deck credibility)
- Zero new setup: Uses llm CLI from Week 1, Groq free tier
- Visceral: Students see the hallucination with their own eyes
- Transferable: The two-prompt pattern applies to any research task
- Complements existing capsules: Pairs with Bias Detective (ethics) and Neural Sandbox (context limits)

Offline / No-Internet Fallback

- Print pre-generated AI responses (one hallucinated, one verified) on cards
- Students play detective: which answer is fabricated?
- Use sticky notes to map out what a "verification layer" looks like on paper
- Role-play: one student is the AI, another is the verifier, third is the founder